

CPO Experiment Plan

Synthetic experiments for the CPO project — implementation guide

CPO project plan

April 27, 2026

Abstract

This document specifies four synthetic experiments that together showcase the mechanism, the scaling behaviour, the operating regime, and the clusters-from-data extension of Cluster-Referenced Preference Optimization (CPO). Each experiment is small enough to run on a laptop in under one minute. For each experiment we state a **goal**, the **exact protocol** (data generation, architecture, training, logging, figure specification), and a **success criterion**: a concrete quantitative prediction that the experiment must either confirm or refute.

Common setup

The four experiments share the following skeleton. Implement once; reuse.

World. Prompts $x \in \{1, \dots, P\}$ with $P = 8$. Responses $y \in \{1, \dots, R\}$ per prompt with $R = 16$. Fixed true latent quality $q(x, y) \in [0, 1]$ drawn uniformly at world-construction time (seeded).

Annotator clusters. $K = 2$ clusters, each with a threshold τ_k and a label-noise rate ϵ_k . An annotator from cluster k who sees sample (x, y) returns label

$$s = \mathbb{I}\{q(x, y) > \tau_k\} \oplus \text{Bernoulli}(\epsilon_k),$$

where $\oplus$ denotes xor (random flip with probability ϵ_k). The cluster prior is $(\pi_A, 1 - \pi_A)$.

Dataset sampling. For experiments A–C: at each training step, sample a mini-batch directly from the population by drawing $(x_i, y_i) \sim \text{Uniform}(P \times R)$, $c_i \sim \text{Cat}(\pi_A, 1 - \pi_A)$, then labelling as above. For experiment D, sample a fixed offline dataset with annotator identity (see §4).

Policy. Tabular cluster-blind softmax: $\pi_\theta(y \mid x) = \text{softmax}(\theta_{x,:})[y]$, $\theta \in \mathbb{R}^{P \times R}$, initialised to zero. Reference $\pi_{\text{ref}}(y \mid x) = 1/R$ uniform.

Log-ratio reward. $r_i = \log \pi_\theta(y_i \mid x_i) - \log \pi_{\text{ref}}(y_i \mid x_i) = \log \pi_\theta(y_i \mid x_i) + \log R$.

Loss. Unary prospect loss:

$$\mathcal{L}(\theta) = \mathbb{E}_i[1 - \sigma(\beta m_i)], \quad m_i = \begin{cases} r_i - z_{c(i)} & s_i = D \\ z_{c(i)} - r_i & s_i = U \end{cases}$$

with $\beta = 1$ unless stated otherwise, and $\lambda_D = \lambda_U = 1$. The reference z is detached in the backward pass (stop-gradient).

Reference point (default variant). Log-ratio-mean over undesirable labels:

$$z_k \approx \mathbb{E}[r_i \mid c_i = k, s_i = U],$$

estimated by an EMA with rate $\rho = 0.1$:

$$z_k \leftarrow (1 - \rho) z_k + \rho \frac{1}{|B_{k,U}|} \sum_{i \in B_{k,U}} r_i,$$

where $B_{k,U}$ is the subset of the mini-batch with cluster k and label U . For KTO, the same statistic is pooled over all clusters.

Optimiser. Adam, learning rate 0.15, batch size 128, 400 training steps.

Evaluation. Expected true quality under the current policy:

$$\mathbb{E}[q] = \frac{1}{P} \sum_x \sum_y \pi_\theta(y \mid x) q(x, y),$$

computed in closed form from the tabular logits. Record every 20 steps.

We pick $\mathbb{E}[q]$ because it measures **what we actually want** (latent response quality), is **independent of the labelling proxy** (so no method gets credit for overfitting to noisy labels), is **cluster-blind** (matching the cluster-blind policy), and is **closed-form computable** (no evaluation noise). It turns the question "did this method extract signal from corrupted labels?" into a single, comparable, well-calibrated number on $[0, 1]$.

Seeds. All experiments use $n_{\text{seeds}} \in \{3, 4, 5\}$ independent seeds. Each seed regenerates the world (true q), the label stream, and the policy initialisation.

1 Experiment A — main comparison with gradient-weight diagnostic

Goal

Demonstrate the core mechanism: CPO outperforms KTO under base-rate heterogeneous annotators because *its per-cluster reference prevents the saturation of the sigmoid for the minority cluster*. The figure must not only show that CPO wins but *why*: per-cluster gradient weights $\sigma'(\beta m_i)$ should look qualitatively different between KTO and CPO.

Protocol

World parameters. $\pi_A = 0.9$. $\tau_A = 0.25$, $\tau_B = 0.75$. $\epsilon_A = \epsilon_B = 0.05$. $n_{\text{seeds}} = 4$.

Methods to compare.

- KTO on the full mixed dataset, single global z .
- CPO on the full mixed dataset, per-cluster z_A, z_B .
- *Oracle-Bob-only*: KTO trained only on the subset with $c_i = 1$ (Bob, strict). This baseline uses 1/10 of the data but clean labels; it is a useful ceiling.

Logging.

- Every 20 steps: evaluate $\mathbb{E}[q]$.
- *For the first seed only*, at every step, record the full per-sample gradient weight $\sigma'(\beta m_i) = \sigma(\beta m_i)(1 - \sigma(\beta m_i))$ together with the cluster ID c_i .

Figure specification

A 1×3 panel figure, total size $\approx 15 \times 4.4$ inches.

Panel (a) — True-quality recovery. x -axis: training step. y -axis: $\mathbb{E}[q]$. Three curves: KTO (solid blue), CPO (solid orange), Oracle-Bob-only (dashed black). Shaded region = ± 1 s.d. over seeds. Legend in lower-right.

Panel (b-KTO) — KTO gradient calibration. x -axis: training step. y -axis: per-cluster *mean* gradient weight $\overline{\sigma'(\beta m_i)_k}$, smoothed with a moving average of window 20. Two curves: Alice (green), Bob (red). Horizontal dotted line at 0.25 (the maximum of σ').

Panel (b-CPO) — CPO gradient calibration. Same as (b-KTO), but for CPO.

Success criterion

- **Quality gap:** $\mathbb{E}[q]_{\text{CPO}} - \mathbb{E}[q]_{\text{KTO}} \geq 0.15$ at the end of training.
- **Oracle beat:** $\mathbb{E}[q]_{\text{CPO}} > \mathbb{E}[q]_{\text{Oracle-Bob-only}}$ at the end of training.
- **Mechanism visible:** In panel (b-KTO), both clusters' gradient weights decay to near zero (< 0.02) by the end of training. In panel (b-CPO), Bob's gradient weight remains elevated (mean over the last 25% of training > 0.02) while Alice's decays normally.

2 Experiment B — the π_A sweep

Goal

Confirm a non-trivial *shape-level* prediction: the advantage $\Delta(\pi_A) = \mathbb{E}[q]_{\text{CPO}} - \mathbb{E}[q]_{\text{KTO}}$ is an inverted-U in π_A , peaking at moderate imbalance and vanishing at both extremes. Matching the predicted shape is much stronger evidence of mechanism than any single-setting comparison.

Protocol

World parameters. Thresholds and noise rates as in Experiment A. Sweep $\pi_A \in \{0.5, 0.7, 0.85, 0.95, 0.99\}$. $n_{\text{seeds}} = 3$.

Methods. KTO and CPO on the full dataset. No oracle in this experiment — the point is the shape of the relative advantage.

Figure specification

A 1×2 panel figure, total size $\approx 11.5 \times 4$ inches.

Panel (left) — Absolute quality. x -axis: π_A . y -axis: final $\mathbb{E}[q]$. Two curves with error bars: KTO (blue circles), CPO (orange squares). Error bars are ± 1 s.d. over seeds.

Panel (right) — Difference. x -axis: π_A . y -axis: $\Delta(\pi_A) = \mathbb{E}[q]_{\text{CPO}} - \mathbb{E}[q]_{\text{KTO}}$ with error bars (propagate error as $\text{s.e.}(\Delta) = \sqrt{\text{s.d.}_{\text{KTO}}^2 + \text{s.d.}_{\text{CPO}}^2} / \sqrt{n_{\text{seeds}}}$). Horizontal black line at $\Delta = 0$.

Success criterion

- **Inverted-U shape:** Δ has a clear maximum at $\pi_A \in \{0.85, 0.95\}$ (not at the endpoints $\pi_A \in \{0.5, 0.99\}$).
 - **Magnitude:** $\max_{\pi_A} \Delta \geq 0.08$.
 - **Endpoints:** $\Delta(\pi_A = 0.5) < 0.04$ and $\Delta(\pi_A = 0.99) < 0.04$.
-

3 Experiment C — four ablations delimiting the regime

Goal

Characterise *when CPO helps and when it does not*. Four independent ablations, each isolating one factor. A good method is one whose failure modes are predictable.

Protocol — C1: reference variant

Setup. Same world as Experiment A ($\pi_A = 0.9$, base-rate heterogeneity). Train KTO and CPO under three reference variants:

1. $r \mid U$ **reference** (default): $z_k = \mathbb{E}[r_i \mid c = k, s = U]$.
2. r **reference:** $z_k = \mathbb{E}[r_i \mid c = k]$ (unconditional mean).
3. **KL reference** (as stated in the CPO paper): $z_k = \mathbb{E}_{x \in C_k} [\text{KL}(\pi_\theta(\cdot \mid x) \parallel \pi_{\text{ref}}(\cdot \mid x))]$.

Success criterion.

- **$r \mid U$ gives a gap:** Under the $r \mid U$ variant, $\text{CPO} \gg \text{KTO}$ (gap ≥ 0.15).
- **r and KL collapse:** Under the r and KL variants, $\text{CPO} \approx \text{KTO}$ (gap < 0.03). *This is a substantive theoretical prediction:* with shared prompts, the KL reference is cluster-invariant because the prompt marginal is cluster-invariant.

Protocol — C2: noise-only heterogeneity

Setup. $\pi_A = 0.9$, but $\tau_A = \tau_B = 0.5$ (same threshold) and $\epsilon_A = 0.3$, $\epsilon_B = 0.02$ (different noise). The clusters differ only in label noise, not in base rate.

Success criterion.

- **No gap:** CPO $\approx$ KTO (gap < 0.03). The CPO mechanism is base-rate calibration, not noise handling.

Protocol — C3: β sweep

Setup. Same world as Experiment A. Sweep $\beta \in \{0.3, 1.0, 3.0, 10.0\}$. $r|U$ reference.

Success criterion.

- **CPO advantage grows with β :** $\Delta(\beta = 10) > \Delta(\beta = 1) > \Delta(\beta = 0.3)$. The saturation-avoidance mechanism is more important at high β .

Protocol — C4: misspecified clusters

Setup. Same heterogeneous world as Experiment A, but the cluster ID passed to CPO is *drawn uniformly at random* (ignoring the true annotator type).

Success criterion.

- **No degradation:** CPO $\approx$ KTO (gap within seed noise, both $\pm$ bars overlap). CPO must not perform worse than KTO when clusters are meaningless.

Figure specification

A 1×4 panel figure, total size $\approx 17 \times 4.2$ inches.

- **Panel C1:** Grouped bar chart with 3 reference-variant groups, each with KTO (blue) and CPO (orange) bars.
 - **Panel C2:** Two-bar chart: KTO, CPO.
 - **Panel C3:** Line plot, x -axis β (log scale), y -axis final $\mathbb{E}[q]$, two curves (KTO, CPO) with error bars.
 - **Panel C4:** Two-bar chart: KTO, CPO.
-

4 Experiment D — learned clusters

Goal

Answer Open Question 2 of the CPO manuscript: can clusters be *learned from data* rather than given? Using a simple annotator-signature-based K -means, we aim to recover the entire oracle benefit. We also demonstrate the bias-variance behaviour over K .

Protocol

Data generation. Static offline dataset with annotator identity.

Algorithm 1 Generate annotator-indexed dataset

- 1: Sample true quality $q(x, y) \sim \text{Uniform}[0, 1]$, shape (P, R) .
 - 2: Draw $N_{\text{ann}} = 120$ annotators, each assigned type $k_a \sim \text{Bernoulli}(1 - \pi_A^*)$ with $\pi_A^* = 0.85$.
 - 3: For each annotator $a = 1, \dots, N_{\text{ann}}$:
 - 4: **for** each of $n_{\text{per}} = 40$ labels **do**
 - 5: Draw $x \sim \text{Uniform}(P)$, $y \sim \text{Uniform}(R)$.
 - 6: Label $s = \mathbb{I}\{q(x, y) > \tau_{k_a}\} \oplus \text{Bernoulli}(0.05)$.
 - 7: Store tuple (x, y, s, a) .
-

Clustering — 1-D signature. For each annotator a , compute the 1-D signature

$$\text{sig}(a) = \frac{1}{|S_a|} \sum_{i \in S_a} s_i = \text{empirical } D\text{-rate of annotator } a,$$

where $S_a = \{i : a_i = a\}$. Run K -means on the $N_{\text{ann}} \times 1$ signature matrix; assign each sample's cluster as $c(i) = \text{cluster}(a_i)$.

Methods to compare at $K = 2$.

- KTO ($K = 1$, single global z).
- CPO with *random* cluster assignment per annotator (NMI ≈ 0 against truth).
- CPO with *learned* cluster assignment (1-D K -means).
- CPO with *oracle* cluster assignment (cheat, use true k_a).

K -sweep. For $K \in \{1, 2, 3, 5, 10\}$, compare:

- CPO with K -means on the 1-D signatures (*learned*).
- CPO with a stratified random partition of annotators into K equal-size groups (*random partition*). This isolates the effect of K on EMA-variance from any effect of clustering quality.

Training. Same CPO loop as in Experiments A–C, but samples are drawn uniformly at random from the pre-generated offline dataset rather than from the online label-generating process. Use the $r|U$ reference. $n_{\text{steps}} = 250$, batch size 128, learning rate 0.2.

Logging. Final $\mathbb{E}[q]$ for each method. For the clustering-quality comparison:

- Normalized mutual information $\text{NMI}(c_{\text{pred}}, k_{\text{true}})$ (use `sklearn.metrics.normalized_mutual_info_score`).
- Purity: $\frac{1}{N_{\text{ann}}} \sum_k \max_{k^*} |\{a : c_{\text{pred}}(a) = k, k_a = k^*\}|$.

Figure specification

A 1×3 panel figure, total size $\approx 15 \times 4.2$ inches.

Panel (a) — Main bar chart. Four bars in order: KTO (blue), CPO random (grey), CPO learned (purple), CPO oracle (orange). y -axis: final $\mathbb{E}[q]$. Error bars = ± 1 s.d. Annotate each bar with its numerical value.

Panel (b) — Cluster recovery. Three groups (random, learned, oracle), each with two bars: NMI (dark) and purity (light). Error bars over seeds.

Panel (c) — K -sweep. x -axis: K (log-scaled). y -axis: final $\mathbb{E}[q]$. Two curves: CPO learned (purple squares), CPO random-partition (grey circles). Two reference horizontal lines: oracle $K = 2$ (orange dashed) and KTO $K = 1$ (blue dotted).

Success criteria

- **Learned $\approx$ oracle:** At $K = 2$, $|\mathbb{E}[q]_{\text{learned}} - \mathbb{E}[q]_{\text{oracle}}| < 0.02$.
 - **Learned $>$ random:** At $K = 2$, $\mathbb{E}[q]_{\text{learned}} - \mathbb{E}[q]_{\text{random}} \geq 0.05$.
 - **Random $\approx$ KTO:** At $K = 2$, $|\mathbb{E}[q]_{\text{random}} - \mathbb{E}[q]_{\text{KTO}}| < 0.03$. Having “multiple clusters” is not enough; the partition must be meaningful.
 - **Clustering recovery:** NMI of learned vs. true annotator type ≥ 0.9 , purity ≥ 0.95 .
 - **Graceful over-clustering:** For $K \in \{3, 5, 10\}$, learned $\mathbb{E}[q]$ stays within 0.02 of the $K = 2$ value.
 - **$K = 1$ collapse:** At $K = 1$, both curves match KTO exactly (no advantage from the CPO machinery when there is only one cluster).
-

5 Reporting deliverables

For each experiment, you should produce:

1. A self-contained Python script that runs the experiment end-to-end on a single machine in under 2 minutes.
2. The figure specified in the protocol, saved as PNG at 150 dpi.
3. A one-paragraph written interpretation addressing each success criterion: did it hold, did it not, and what does that tell us?
4. Cached intermediate results (e.g., a `.pkl` file) so the figure can be re-plotted without re-running the experiment.

6 Suggested order and timeline

- **Task 1.** Implement the common setup (world, policy, training loop, evaluation). Verify the loop reproduces KTO under $K = 1$ and zero heterogeneity.
- **Tasks 2–3.** Experiment A. The gradient-weight diagnostic is the most intellectually important part of the entire plan; take the time to stare at it.

- **Task 4.** Experiment B (fast; the sweep is cheap).
- **Tasks 5–6.** Experiment C. Four ablations; each one is its own mini-study.
- **Tasks 7–8.** Experiment D. Introduces scikit-learn for clustering.